Model Compression with Two-stage Multi-teacher Knowledge Distillation for Web Question Answering System

Paper Link: https://arxiv.org/abs/1910.08381

Contents:

The Problem Statement

- Deep pre-training and fine-tuning models suffer from **slow inference speed** due to the sheer amount of model parameters
- Applying these complex models to real business scenarios is a challenging but practical problem.
- Model compression methods suffer from **information loss**, leading to inferior models.

Proposed Solution

Two-stage Multi-teacher Knowledge Distillation method for web Question Answering system

- General Q&A distillation task for student model pre-training
- Further fine-tune pretrained student model with multi-teacher knowledge distillation on downstream tasks (MNLI, SNLI, RTE tasks from GLUE).
- Effectively reduces the overfitting bias in individual teacher models and transfers more general knowledge to the student model
- This method significantly achieves comparable results with the original teacher models, along with substantial speedup of model inference.

Scope of the project

- Decide teacher
- Create student by experimentation
- Implement 1-o-1 model
- Implement m-o-m model
- Implement m-o-1 model
  - Stage 1: 3 BERT Teacher models, 1 student model assuming RTE to be the large corpus
  - Stage 2: Fine-Tune student model on other datasets

Knowledge Distillation

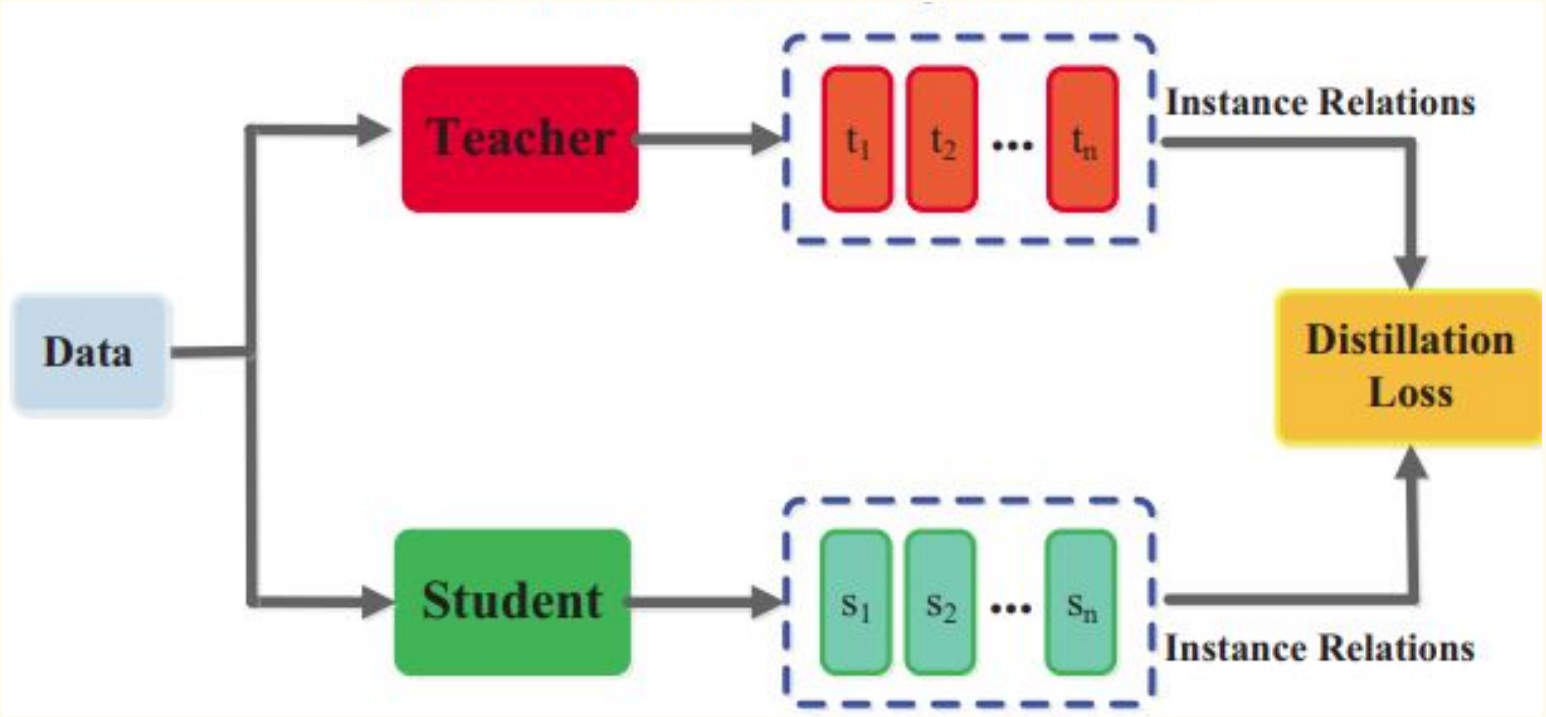

Approaches to Knowledge Distillation

- 1-o-1 model
- m-o-m ensemble model
- m-o-1 model

Question Answering Relevance

Question:	*What can I do when I have headache?*
Passage:	*Drinking warm water mixed with juice squeezed from one-half of a lemon will reduce the intensity of a headache. This particular remedy is beneficial for headaches caused by gas in the stomach. Another option is to apply lemon crusts, pounded into a paste, on your forehead to immediately relieve pain...*
Label:	*Relevant*

The Two-stage Multi-teacher Knowledge Distillation Approach

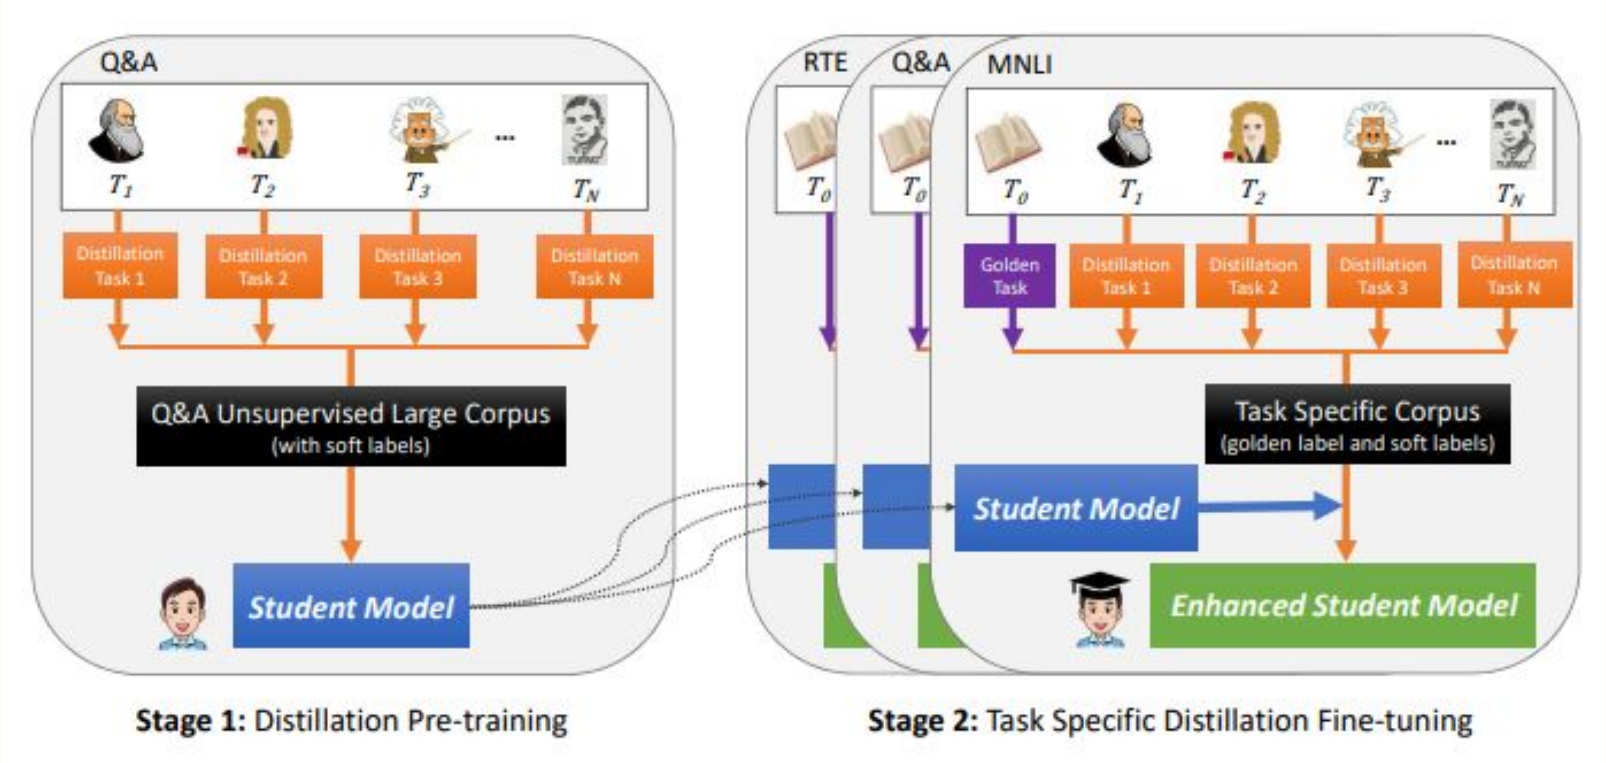

Datasets Used

- **MNLI - Multi-Genre Natural Language Inference**
  - A collection of paired sentences labeled as entailment, contradiction, or neutral
  - Used for natural language inference tasks.
- **SNLI - Stanford Natural Language Inference**
  - A collection of paired sentences labeled as entailment, contradiction, or neutral
  - Used for natural language inference tasks.
- **QNLI - Question-answering Natural Language Inference**
  - A collection of paired questions and sentences labeled as entailment or not entailment
  - Used for natural language inference tasks.
- **RTE - Recognizing Textual Entailment**
  - A collection of paired sentences labeled as entailment or not entailment
  - Used for natural language inference tasks, with a focus on recognizing textual entailment in natural language.

Our implementation

- Prerequisite experimentation *(Results shown later)*
  - Implementation of 1-o-1 KD model on MNLI dataset
  - Implemented m-o-m KD model using ensemble method
  - Trained and tested ensemble model on different datasets
- m-o-1 multi-teacher knowledge distillation to one student assuming RTE as large corpus (*first stage*)
- Fine-tuning the student model on specific task using QNLI (*second stage*)

Concept used

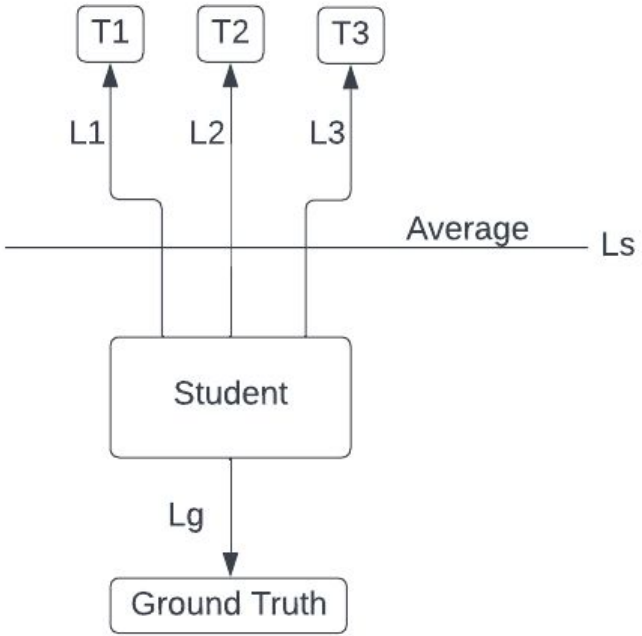

Student loss $= (1 - \alpha) l_g + \alpha * l_s$

Where α is a loss weight ratio

Experiments conducted

- Student model consists of first layer as BERT model layer, and experiments done by increasing the number of subsequent fully connected layers
    - 1 layer
    - 2 layers
    - 3 layers
- Optimizer chosen between ADAM and SGD
- Created 1-o-1 KD model on MNLI, QNLI, SNLI and RTE
- Created m-o-m ensemble model by making multiple teachers and corresponding 1 layer student, and voting logic for final student

Results

- Accuracies of experiments on student model (1 bert layer + 1 fc layer) for deciding the optimizer:
  - Teacher model : Bert base uncased
  - Dataset : MNLI

	Train Accuracy	Validation Accuracy	Test Accuracy
Adam Optimizer	0.997	0.36817	0.39545
SGD Optimizer	0.962	0.2929	0.3163

Results

- Accuracies of experiments on student model (Adam optimizer) for deciding the number of layers:
  - Teacher model : Bert base uncased
  - Dataset : MNLI

	Train Accuracy	Validation Accuracy	Test Accuracy
1 bert layer + 1 fc layer	0.997	0.36817	0.39545
1 bert layer + 2 fc layers	0.996	0.37727	0.38183
1 bert layer + 3 fc layers	0.986	0.36817	0.36817

Results

- Accuracies of teacher (bert base uncased) - student (Adam optimizer and 1 bert layer + 1 fc layer) model experiments on different datasets:
  - Teacher 1 (T1) : Epochs = 5, Learning rate = 2e-5,
  - Teacher 2 (T2) : Epochs = 5, Learning rate = 3e-5
  - Teacher 3 (T3) : Epochs = 5, Learning rate = 5e-5
  - Student 1 (S1)
  - Student 2 (S2)
  - Student 3 (S3)
  - Student (S) : Majority voting ensemble

Results

MNLI	T1	0.5182	T2	0.6681	T3	0.5818		
	S1	0.4275	S2	0.3727	S3	0.3772	S	0.3681
SNLI	T1	0.54	T2	0.50	T3	0.6428		
	S1	0.35	S2	0.438	S3	0.520	S	0.469
QNLI	T1	0.7769	T2	0.8417	T3	0.8417		
	S1	0.5755	S2	0.4748	S3	0.5755	S	0.5467
RTE	T1	0.568	T2	0.589	T3	0.517		
	S1	0.625	S2	0.589	S3	0.611	S	0.5948

Results

- Accuracy of m-o-1 model

Creating student model

```
student_11 = MLP(0.1).to(device)
student_11_optim = torch.optim.Adam(student_11.parameters(), lr=1e-4)
train_logits_list_1 = [train_logits_11, train_logits_12, train_logits_13]
val_logits_list_1 = [val_logits_11, val_logits_12, val_logits_13]
test_logits_list_1 = [test_logits_11, test_logits_12, test_logits_13]
train_pred_labels_student_11, val_pred_labels_student_11, test_pred_labels_student_11, train_pred
```

```
Training..........
Checkpoint accessing..........
Resuming training from epoch 21
Epoch 30   Train loss: 0.0290820110142231    Train accuracy: 0.962
Epoch 40   Train loss: 0.02895363214612007    Train accuracy: 0.97
Evaluating on training data.................
1000 1
Training loss: 0.02603114864230156 Train accuracy:  0.969
Evaluating on validation data.................
138 1
Validation loss: 0.04896077340927677 Validation accuracy:  0.5869565217391305
Evaluating on testing data.................
139 1
Testing loss: 0.04710530548644581 Test accuracy:  0.5827338129496403
```

Results

- Accuracy of m-o-1 model

Fine tuning pre-trained student on qnli dataset

```
train_logits_list_2 = [train_logits_21, train_logits_22, train_logits_23]
val_logits_list_2 = [val_logits_21, val_logits_22, val_logits_23]
test_logits_list_2 = [test_logits_21, test_logits_22, test_logits_23]
train_pred_labels_student_12, val_pred_labels_student_12, test_pred_labels_student_12, train_pred
```

```
Training..........
No saved checkpoints to resume
Epoch 0   Train loss: 0.04956310951709747    Train accuracy: 0.509
Epoch 10   Train loss: 0.028006032228469847    Train accuracy: 0.975
Epoch 20   Train loss: 0.027910725146532058    Train accuracy: 0.977
Epoch 30   Train loss: 0.027913129150867463    Train accuracy: 0.978
Epoch 40   Train loss: 0.027914887696504593    Train accuracy: 0.978
Evaluating on training data.................
1000 1
Training loss: 0.025544155269861223 Train accuracy:  0.978
Evaluating on validation data.................
220 1
Validation loss: 0.05101537812839855 Validation accuracy:  0.4818181818181818
Evaluating on testing data.................
220 1
Testing loss: 0.04940701127052307 Test accuracy:  0.5318181818181819
```

Comparison

- Comparison between 1-o-1, m-o-m and m-o-1 model

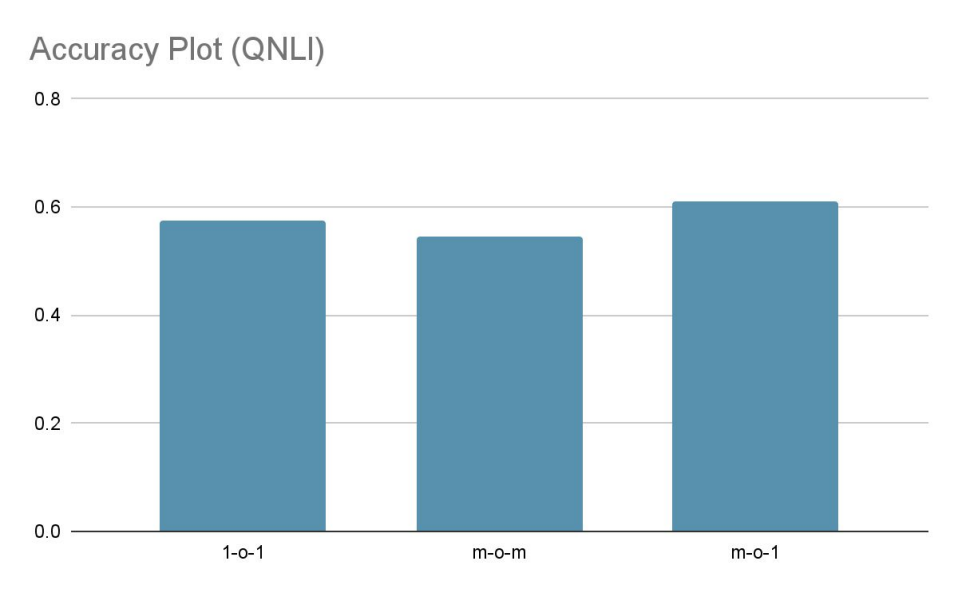

Challenges faced

- In the first stage, a large corpus of Q&A dataset was derived from unlabelled data obtained from commercial search engine using BERT Large. We did not have access to this dataset hence we had to assume RTE as our large corpus (hence the lower accuracies)
- Limited resources with respect to computation power (google colab gpu runtime kept crashing)
- Understanding the loss function of the first stage of multi-teacher knowledge distillation

Thank You!